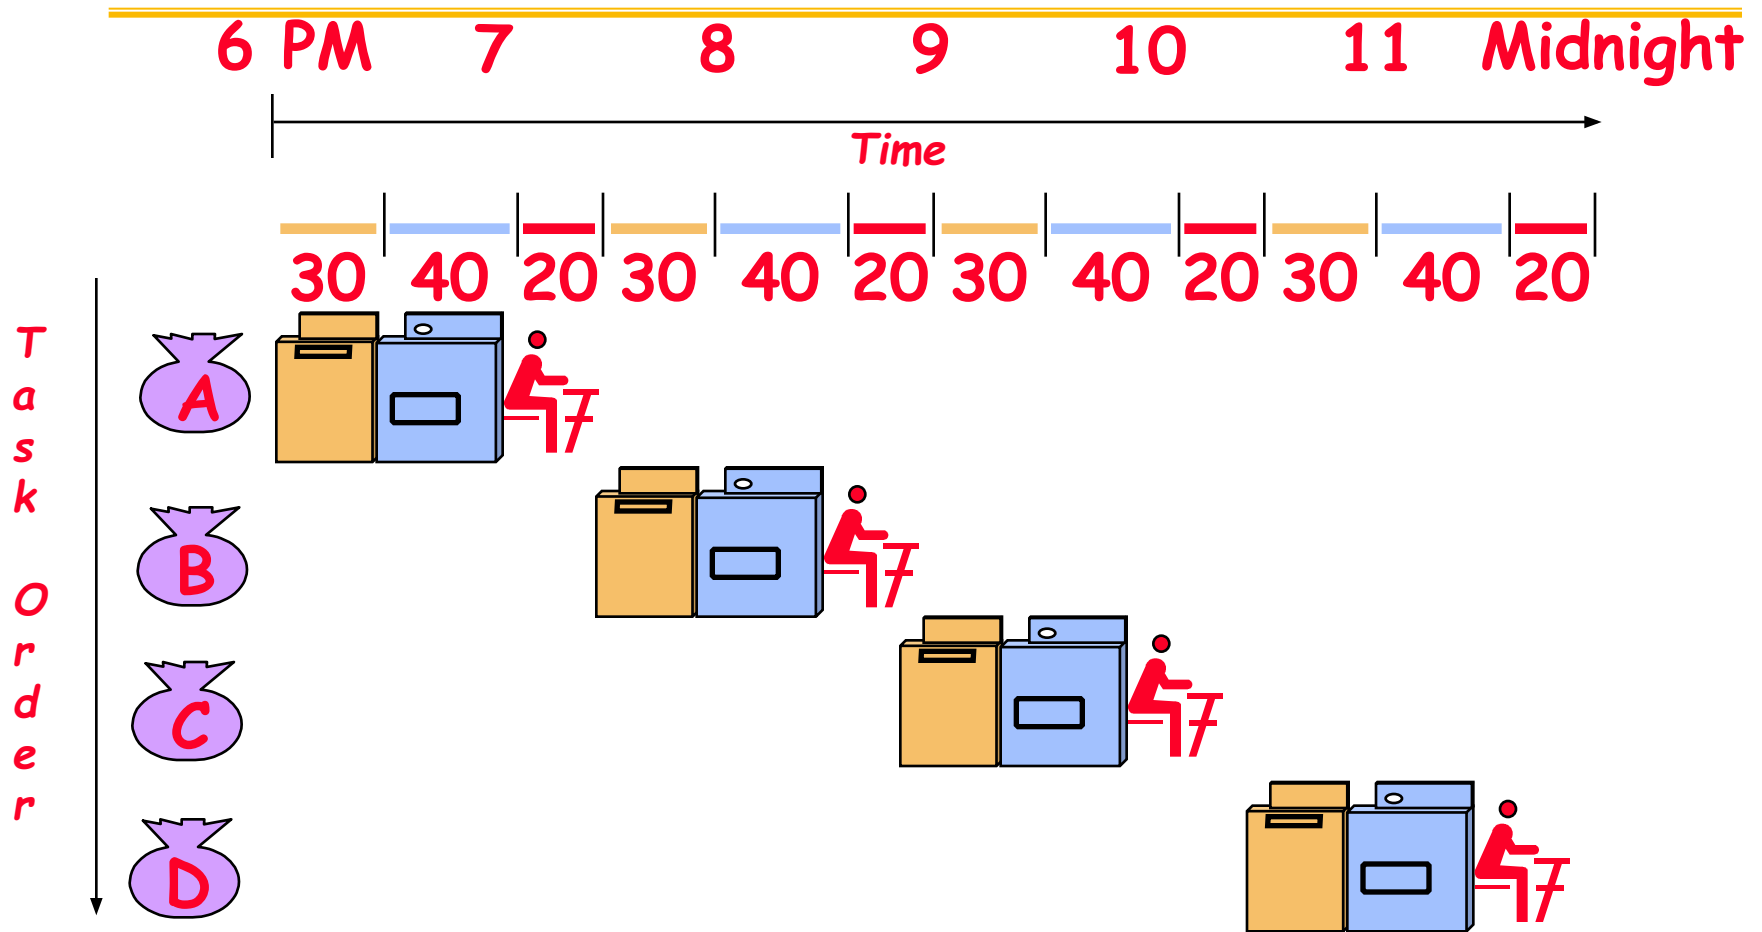


Pipelining

Outline

- **5 stage pipelining**
- **Structural and Data Hazards**
- **Forwarding**
- **Branch Schemes**
- **Multi-cycle Operations**
- **Conclusion**

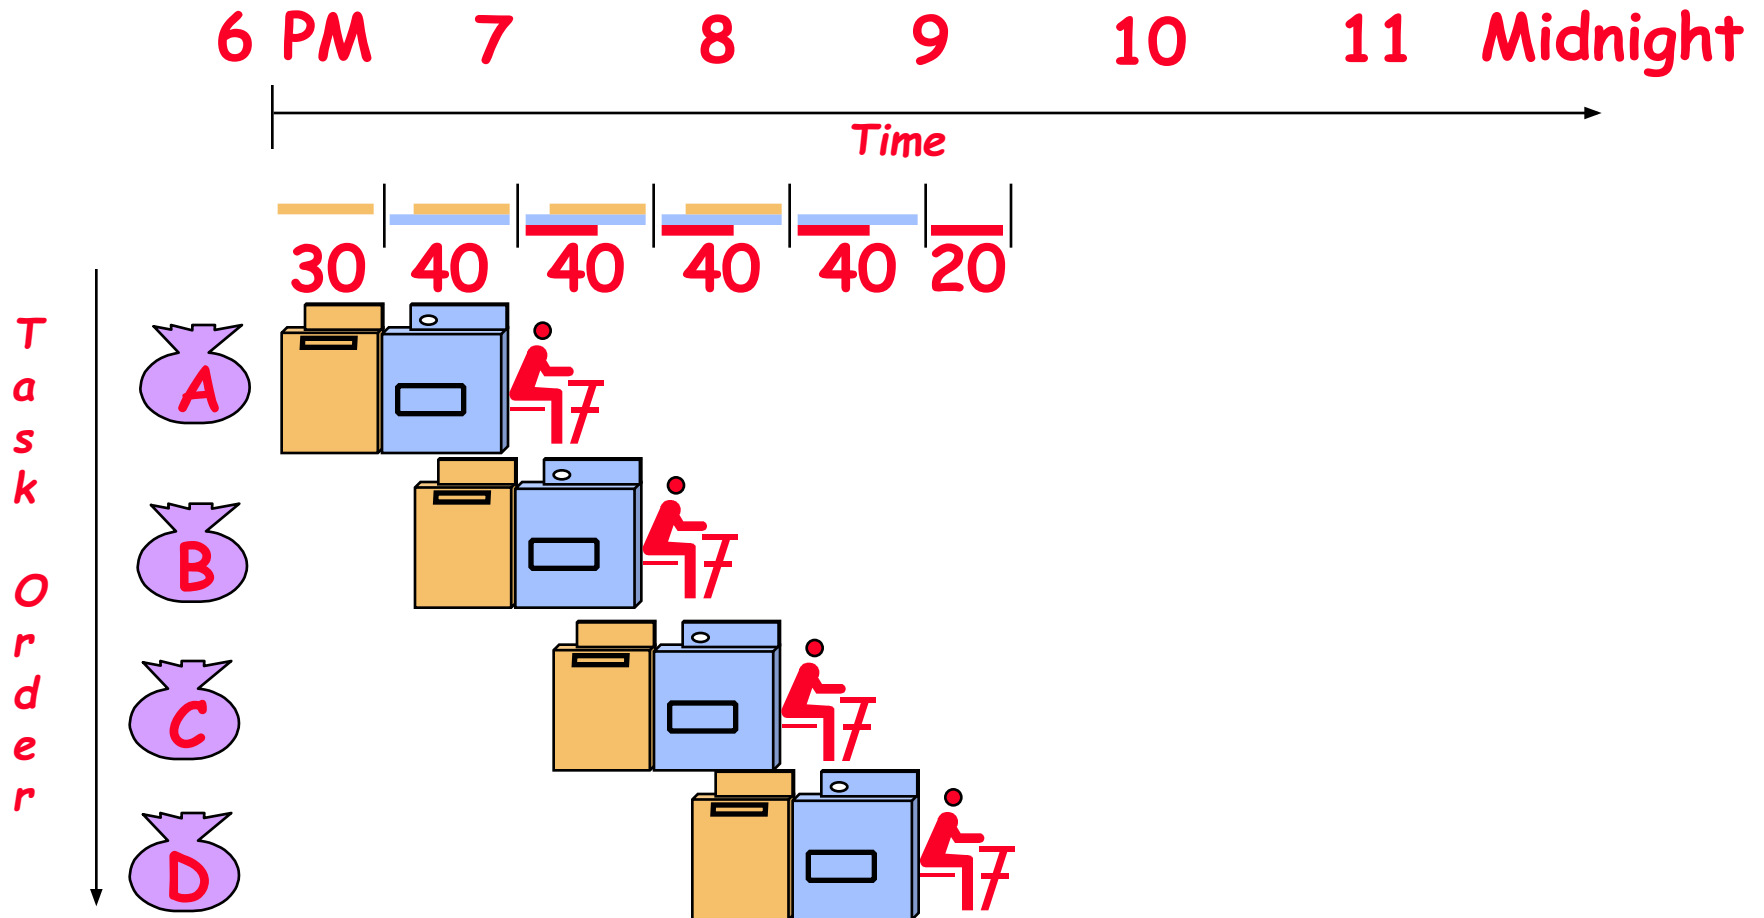
Sequential Laundry



- Sequential laundry takes 6 hours for 4 loads
- If they learned pipelining, how long would laundry take?


Pipelined Laundry

Start work ASAP



- Pipelined laundry takes 3.5 hours for 4 loads

Pipelining Lessons




- Pipelining doesn't help **latency** of single task, it helps **throughput** of entire workload
- Pipeline rate limited by **slowest** pipeline stage
- **Multiple** tasks operating simultaneously
- Potential speedup = **Number pipe stages**
- Unbalanced lengths of pipe stages reduces speedup
- Time to "**fill**" pipeline and time to "**drain**" it reduces speedup

Instruction Pipelining


- Execute billions of instructions, so *throughput* is what matters
 - except when?
- What is desirable in instruction sets for pipelining?
 - Variable length instructions vs. all instructions same length?
 - Memory operands part of any operation vs. memory operands only in loads or stores?
 - Register operand many places in instruction format vs. registers located in same place?

Example: MIPS (Note register location)

Register-Register



Register-Immediate



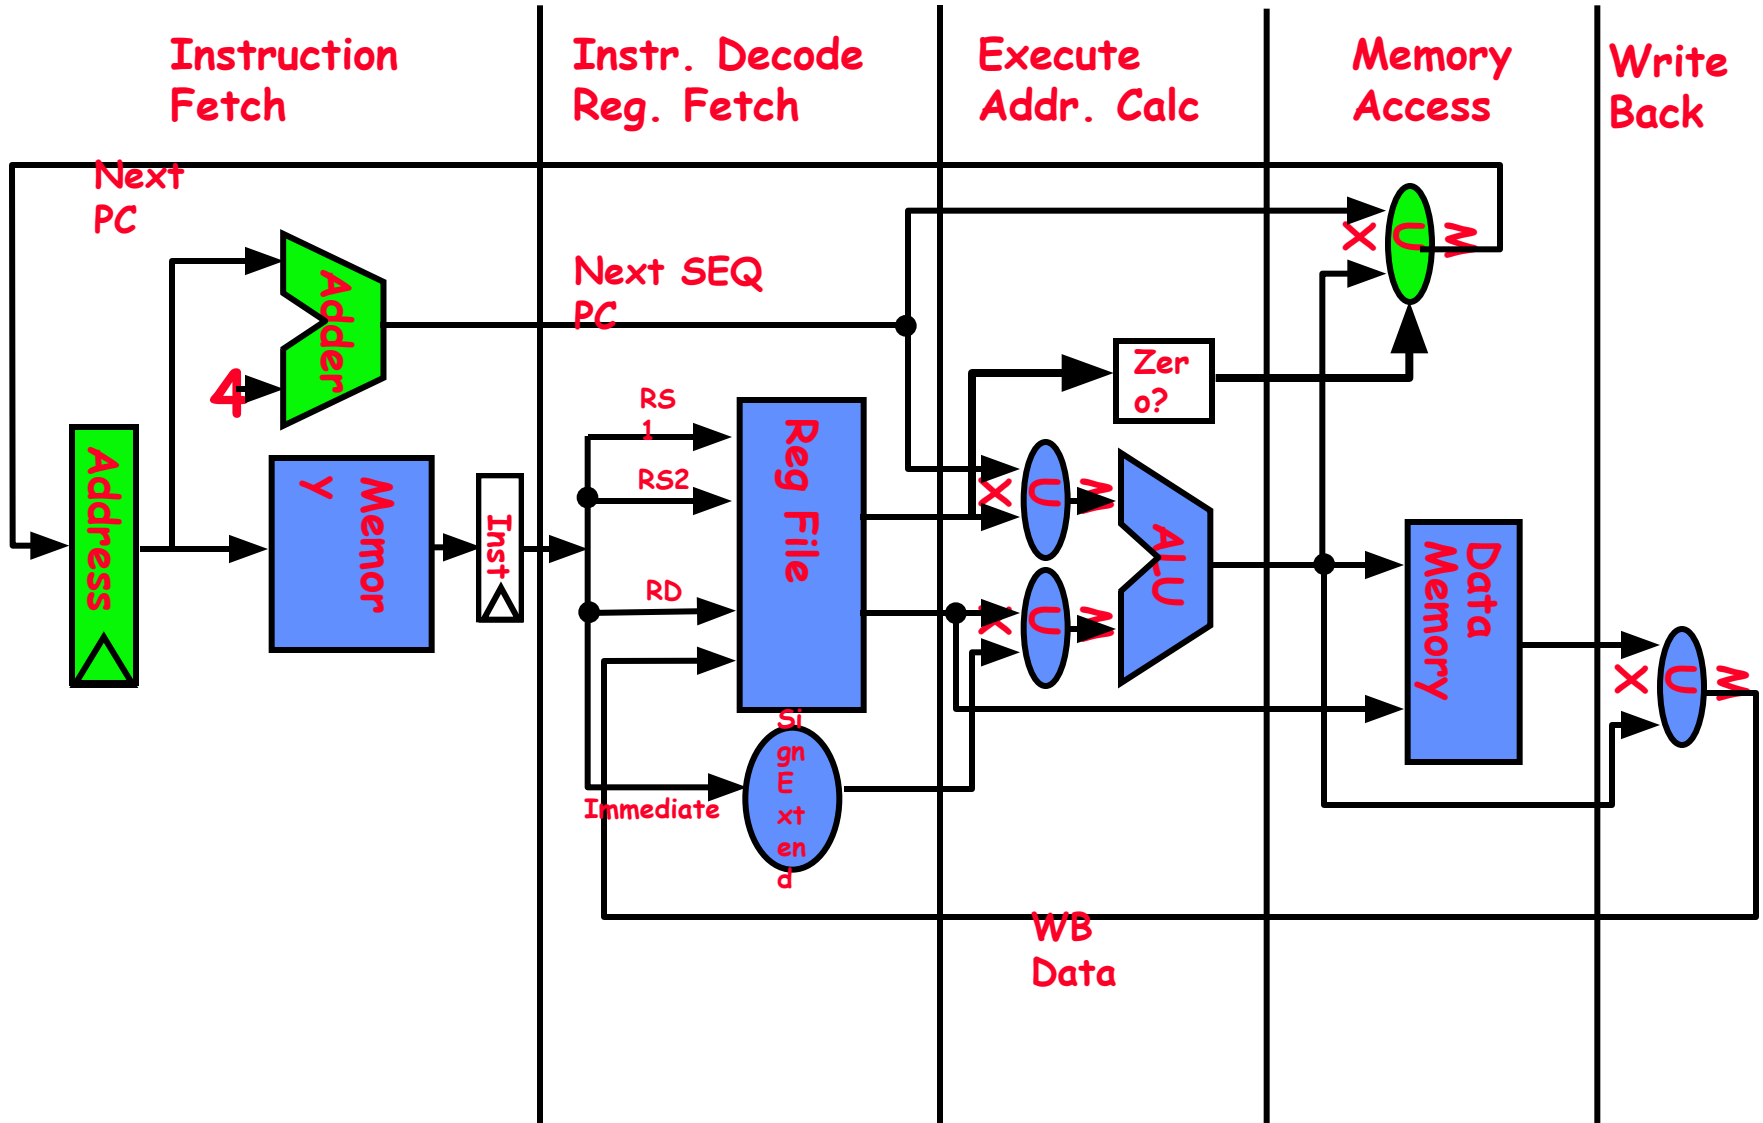
Branch



Jump / Call

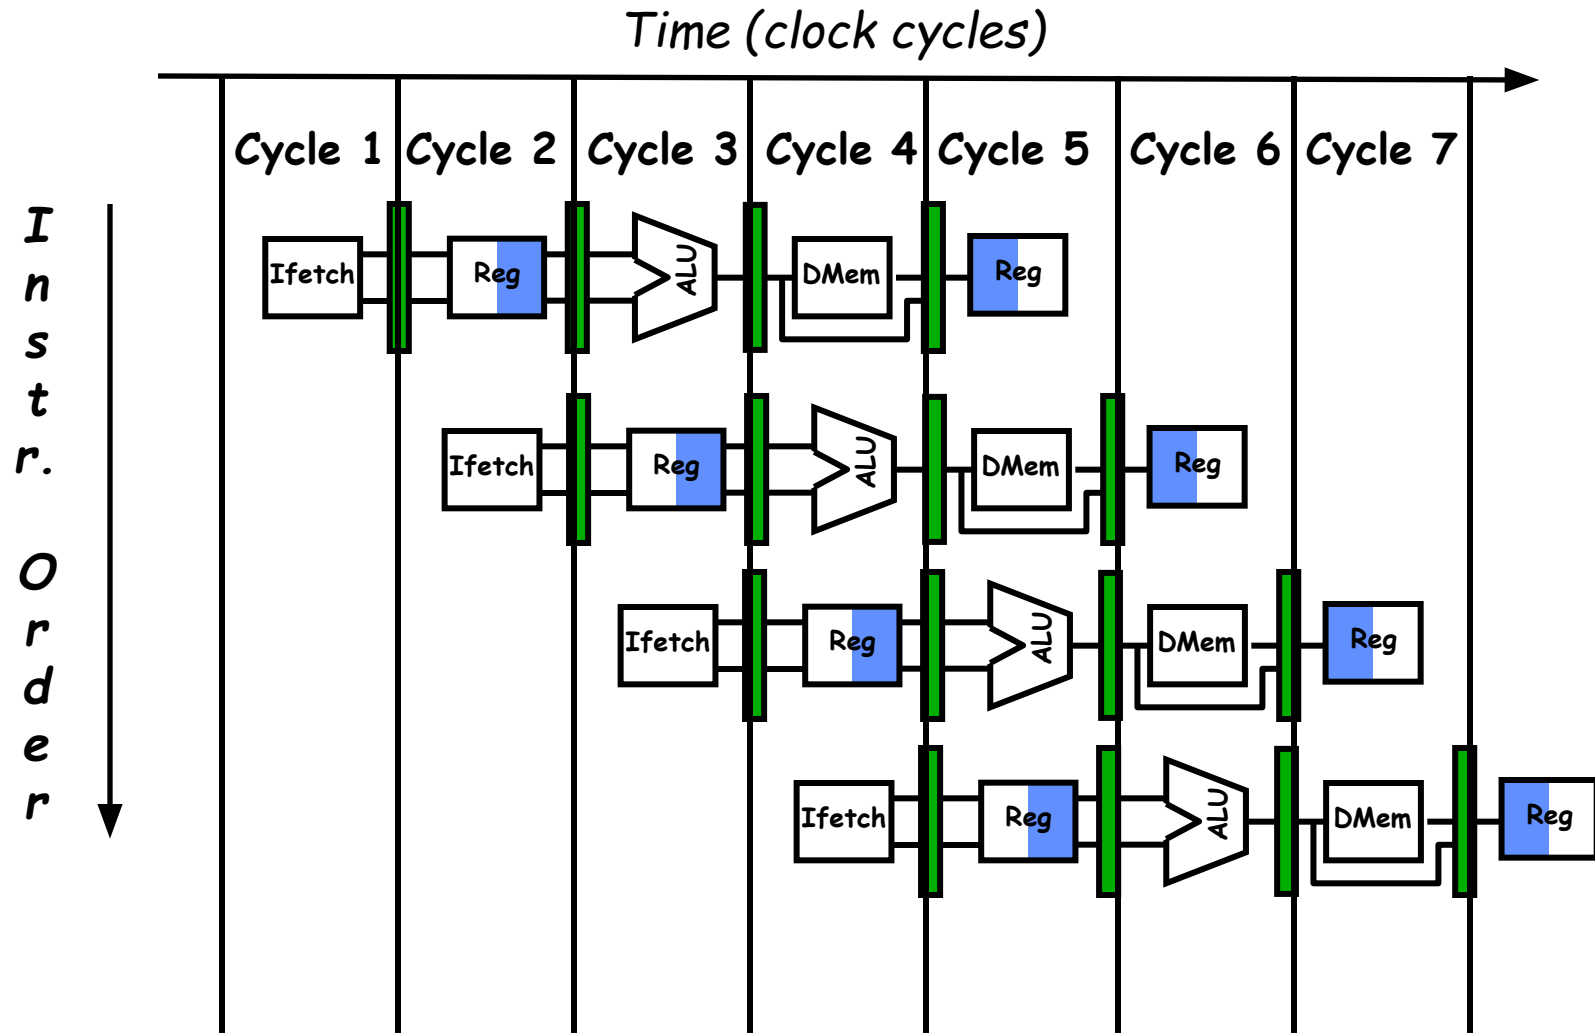


5 Steps of MIPS Datapath



Five-Stage Pipelining

Figure A.2, Page A-8

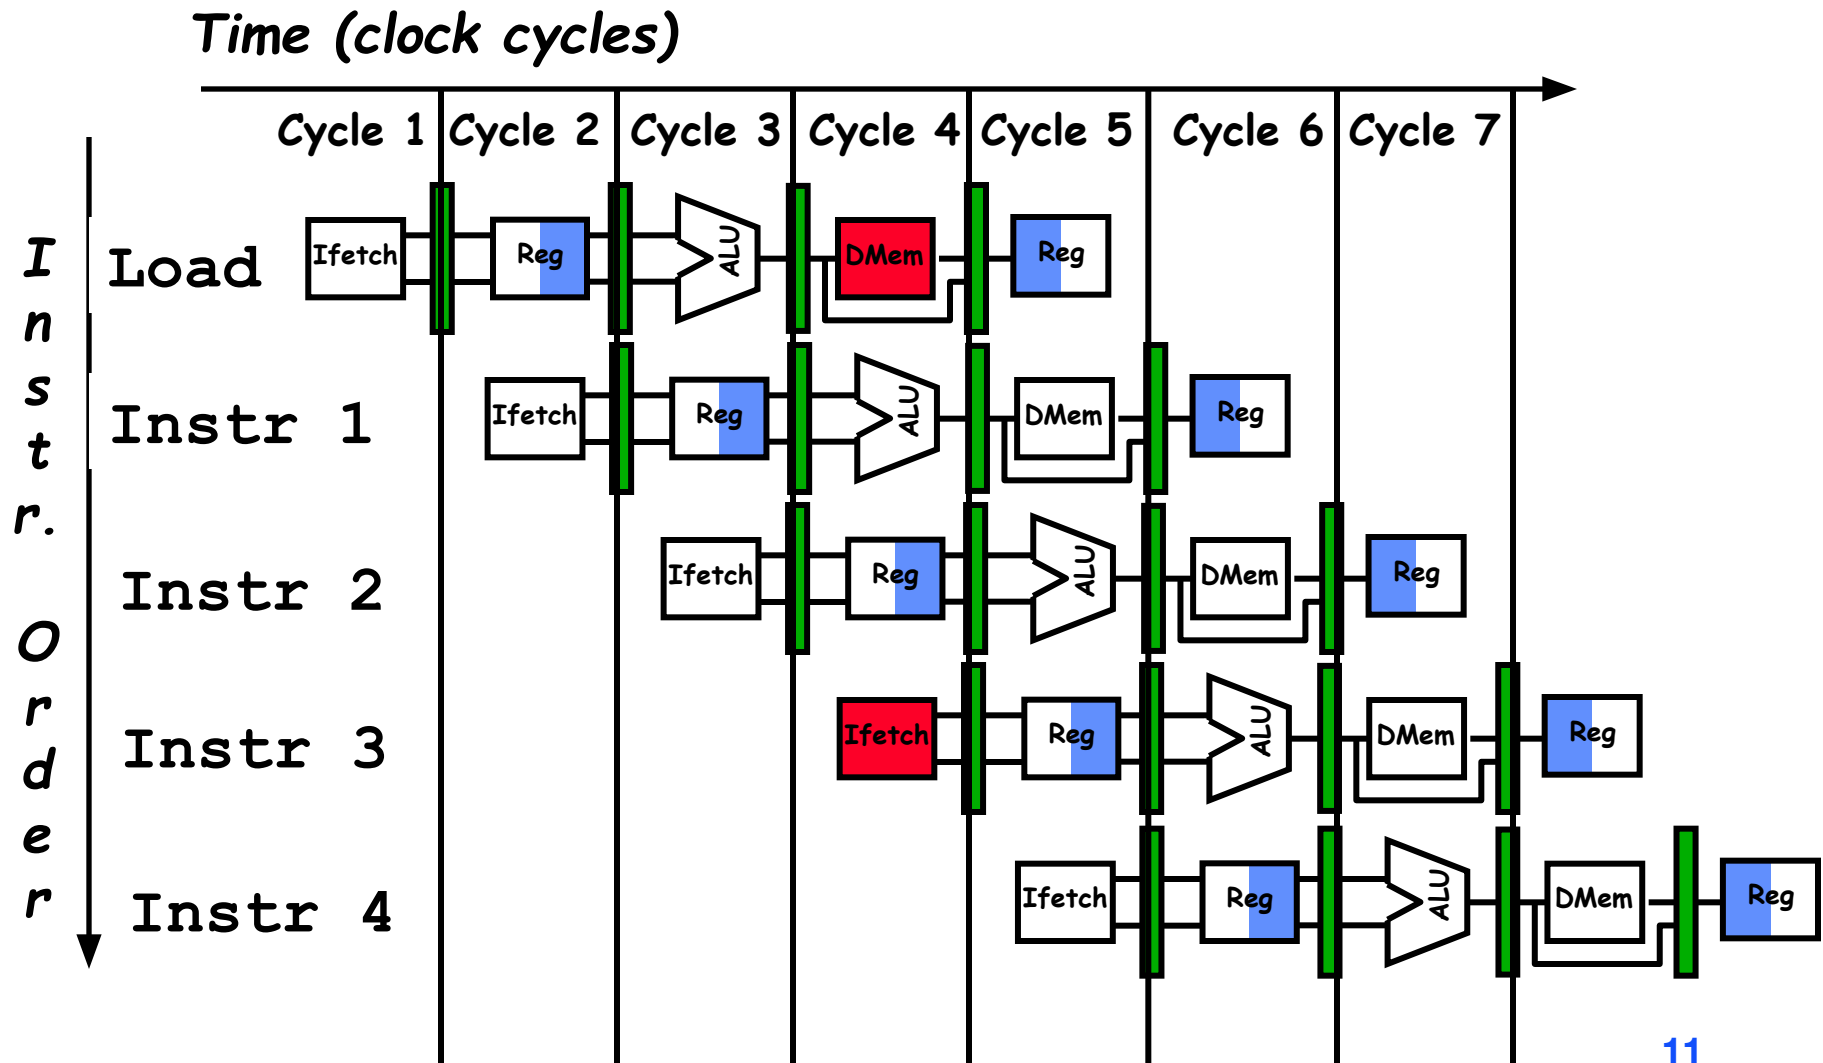


Pipelining is not quite that easy!

- Limits to pipelining: **Hazards** prevent next instruction from executing during its designated clock cycle
 - **Structural hazards**: HW cannot support this combination of instructions (single person to fold and put clothes away)
 - **Data hazards**: Instruction depends on result of prior instruction still in the pipeline (missing sock)
 - **Control hazards**: Caused by delay between the fetching of instructions and decisions about changes in control flow (branches and jumps).

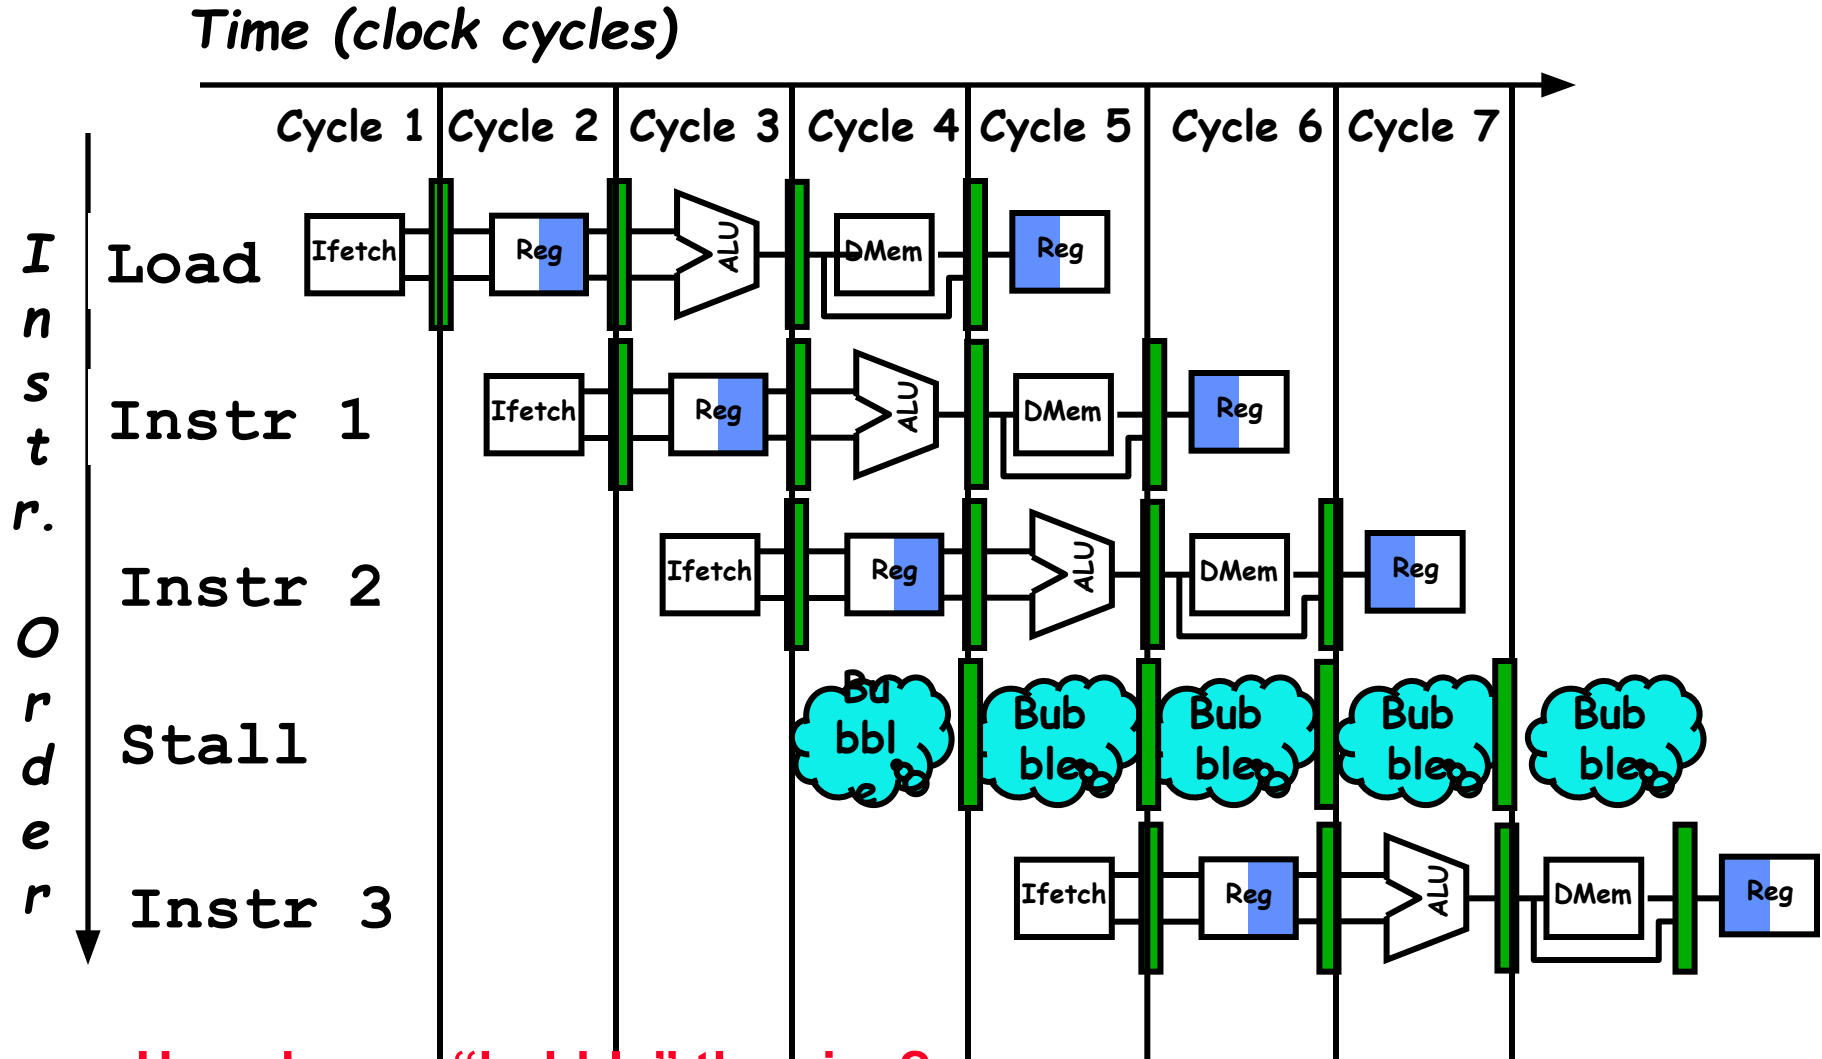
One Memory Port/Structural Hazards

Figure A.4, Page A-14



One Memory Port/Structural Hazards

(Similar to Figure A.5, Page A-15)

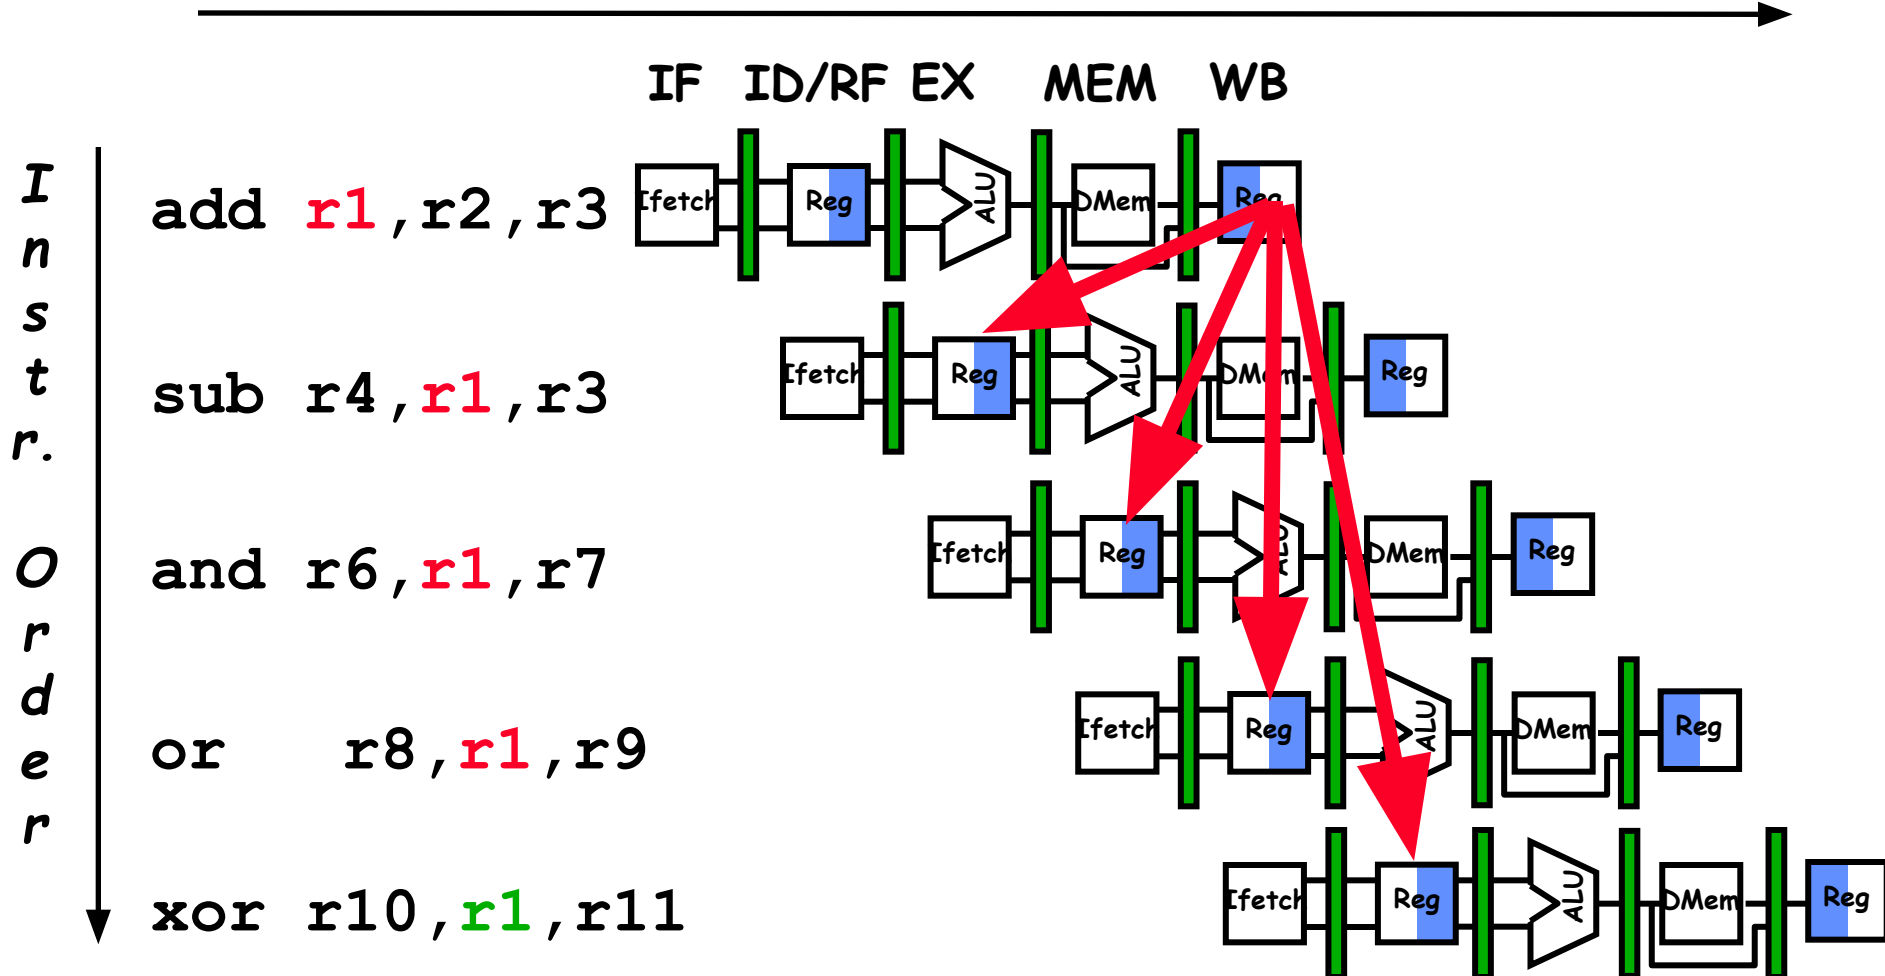


How do you “bubble” the pipe?

Data Hazard on R1

Figure A.6, Page A-17

Time (clock cycles)



Three Generic Data Hazards

- **Read After Write (RAW)**

Instr_j tries to read operand before Instr_i writes it

$\left(\begin{array}{l} \text{I: add } \textcolor{red}{r1}, r2, r3 \\ \text{J: sub } r4, \textcolor{red}{r1}, r3 \end{array} \right.$

- Caused by a “**Dependence**” (in compiler nomenclature). This hazard results from an actual need for communication.

Three Generic Data Hazards

- **Write After Read (WAR)**

Instr_j writes operand before Instr_i reads it


```
    I:  sub  r4, r1, r3
    J:  add  r1, r2, r3
    K:  mul  r6, r1, r7
```

- Called an “**anti-dependence**” by compiler writers. This results from reuse of the name “**r1**”.
- Can’t happen in MIPS 5 stage pipeline because:
 - All instructions take 5 stages, and
 - Reads are always in stage 2, and
 - Writes are always in stage 5

Three Generic Data Hazards

- **Write After Write (WAW)**

Instr_j writes operand before Instr_i writes it.

 I: sub **r1**, r4, r3
J: add **r1**, r2, r3
K: mul r6, r1, r7

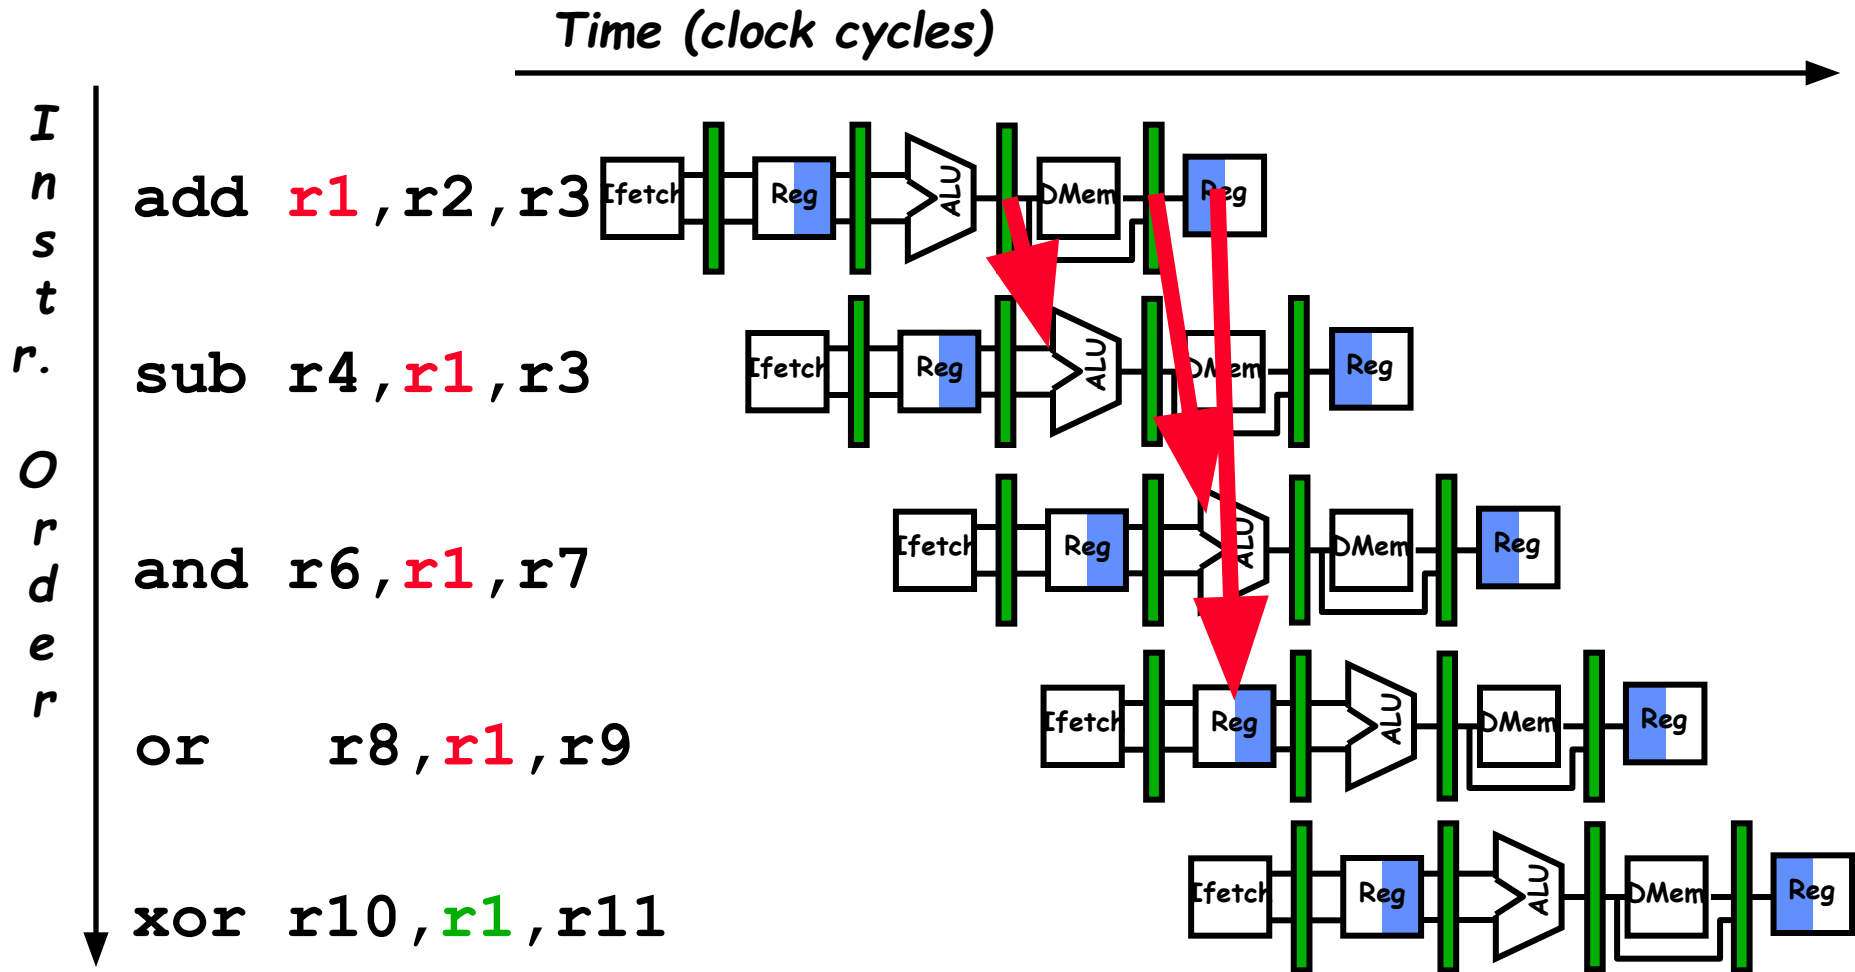
- Called an “**output dependence**” by compiler writers
This also results from the reuse of name “**r1**”.
- Can’t happen in MIPS 5 stage pipeline because:
 - All instructions take 5 stages, and
 - Writes are always in stage 5
- Will see WAR and WAW in more complicated pipes

Outline

- 5 stage pipelining
- Structural and Data Hazards
- **Forwarding**
- **Branch Schemes**
- **Multi-cycle Operations**
- **Conclusion**

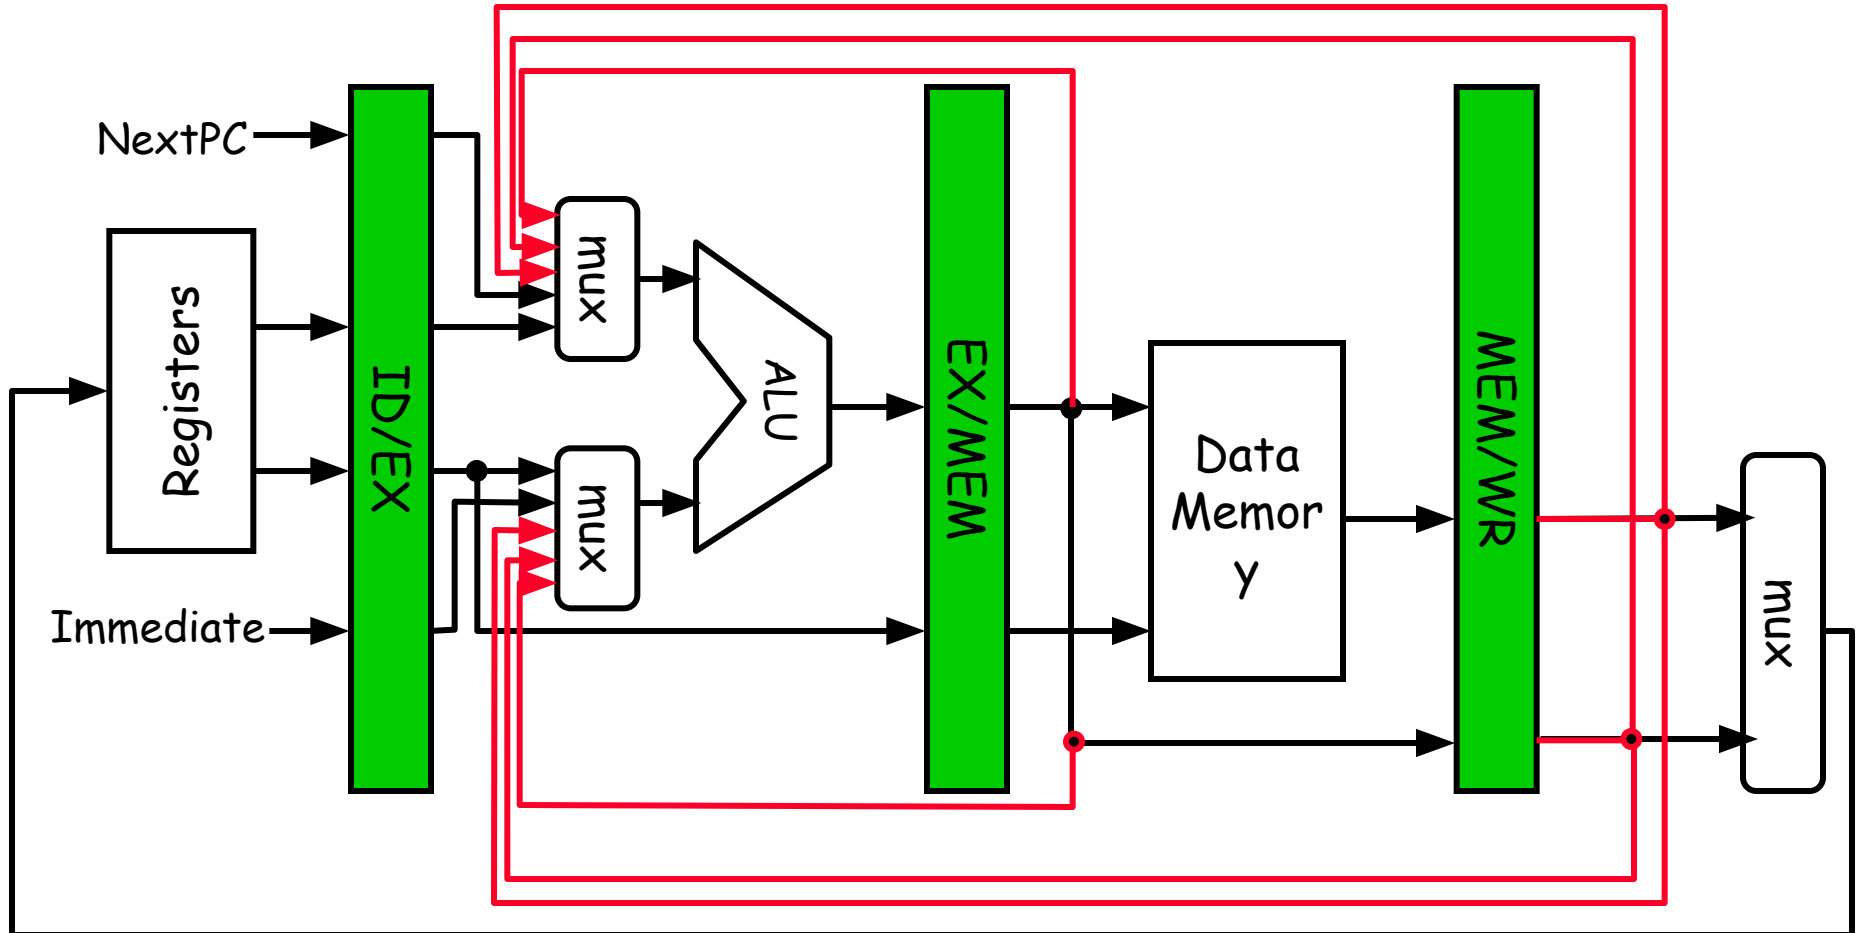
Forwarding to Avoid Data Hazard

Figure A.7, Page A-19



HW Change for Forwarding

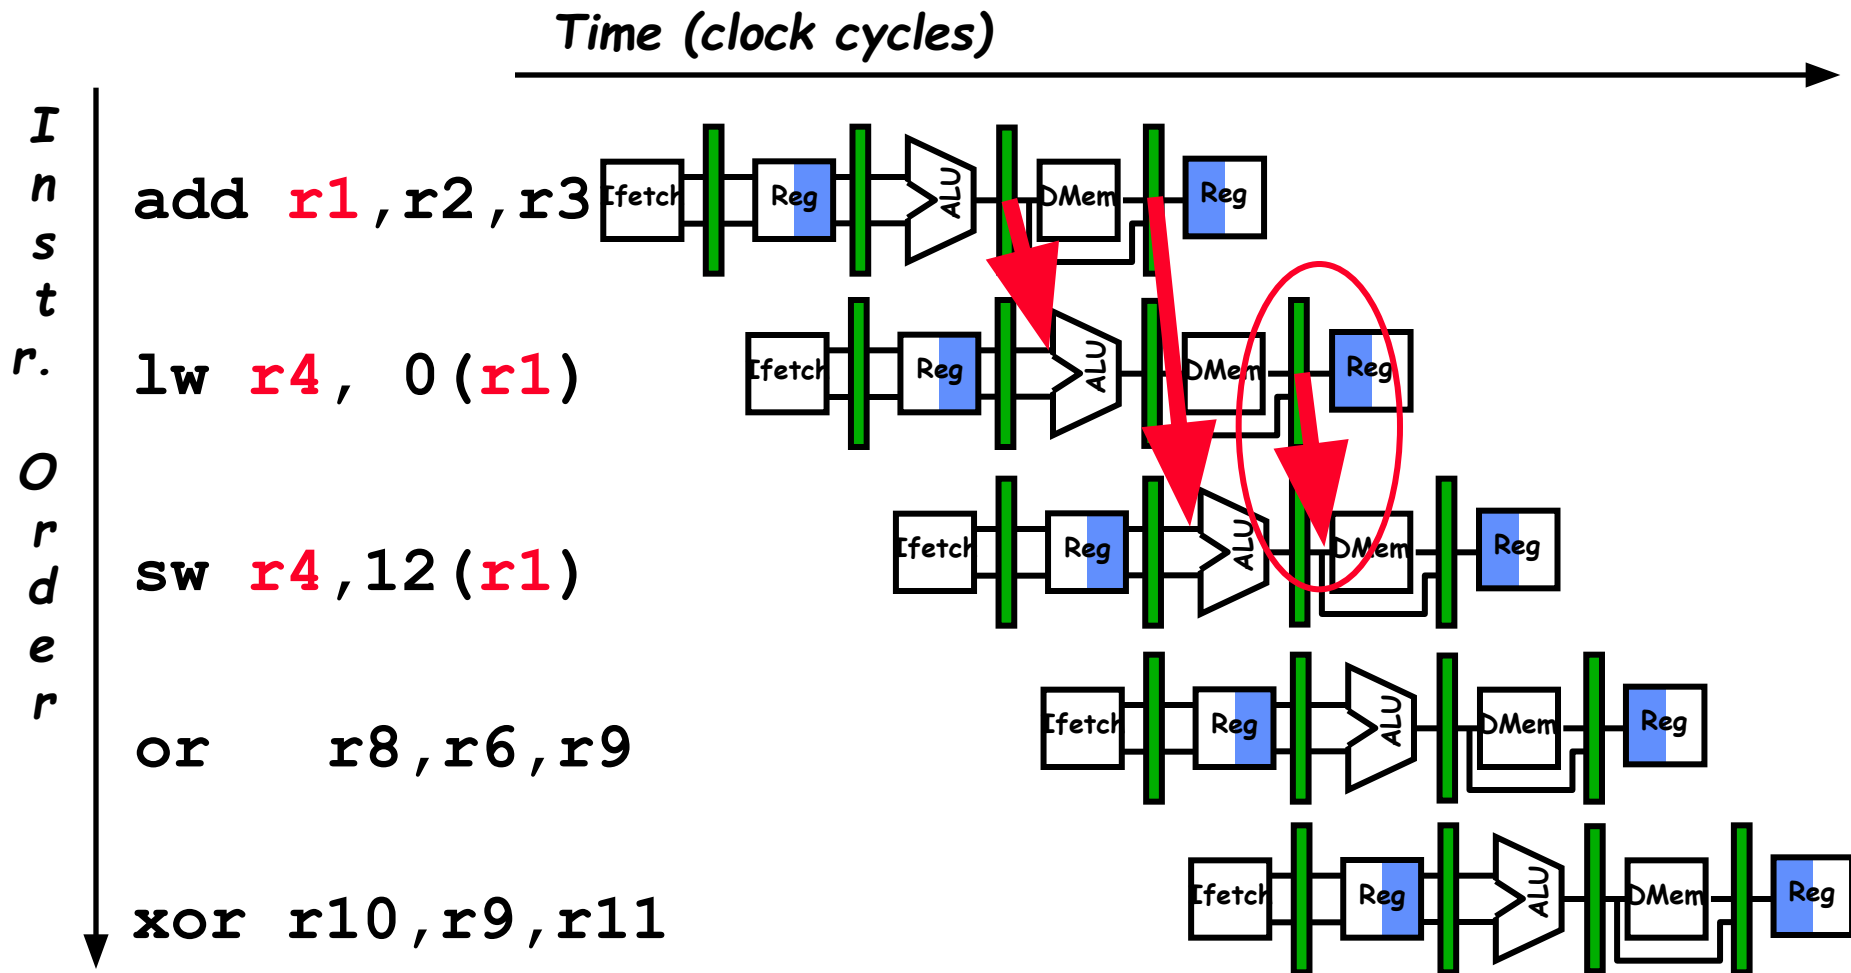
Figure A.23, Page A-37



What circuit detects and resolves this hazard?

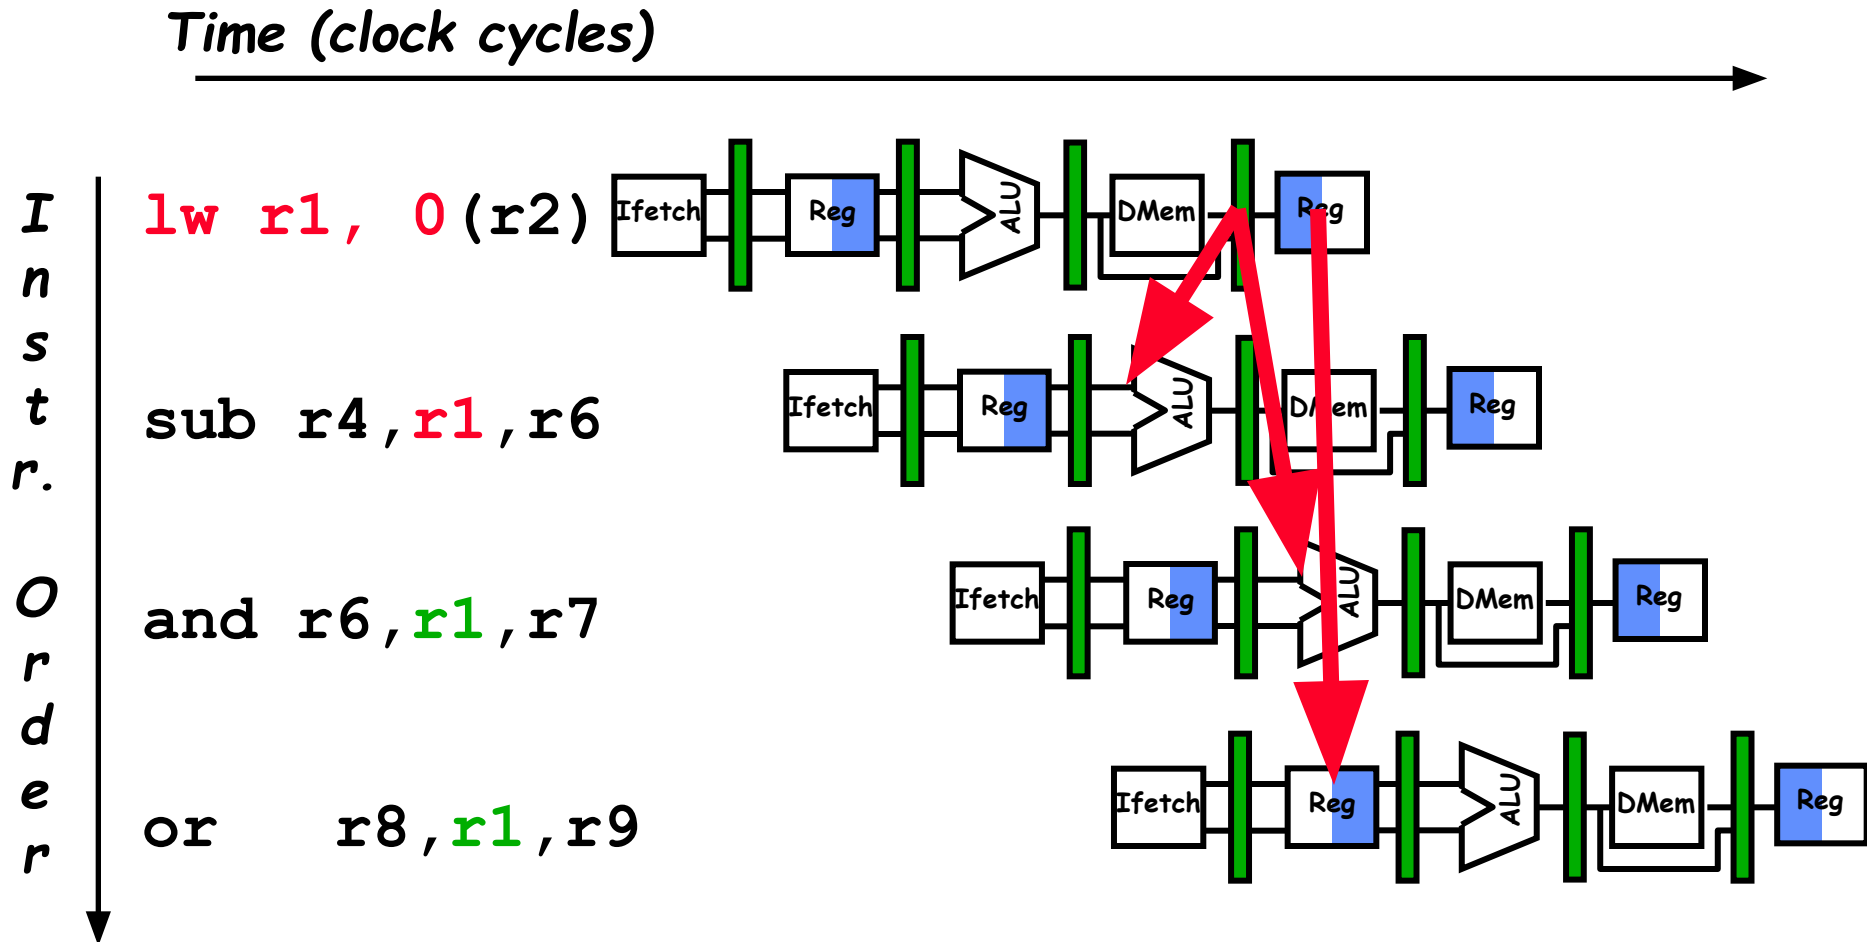
Forwarding to Avoid LW-SW Data Hazard

Figure A.8, Page A-20



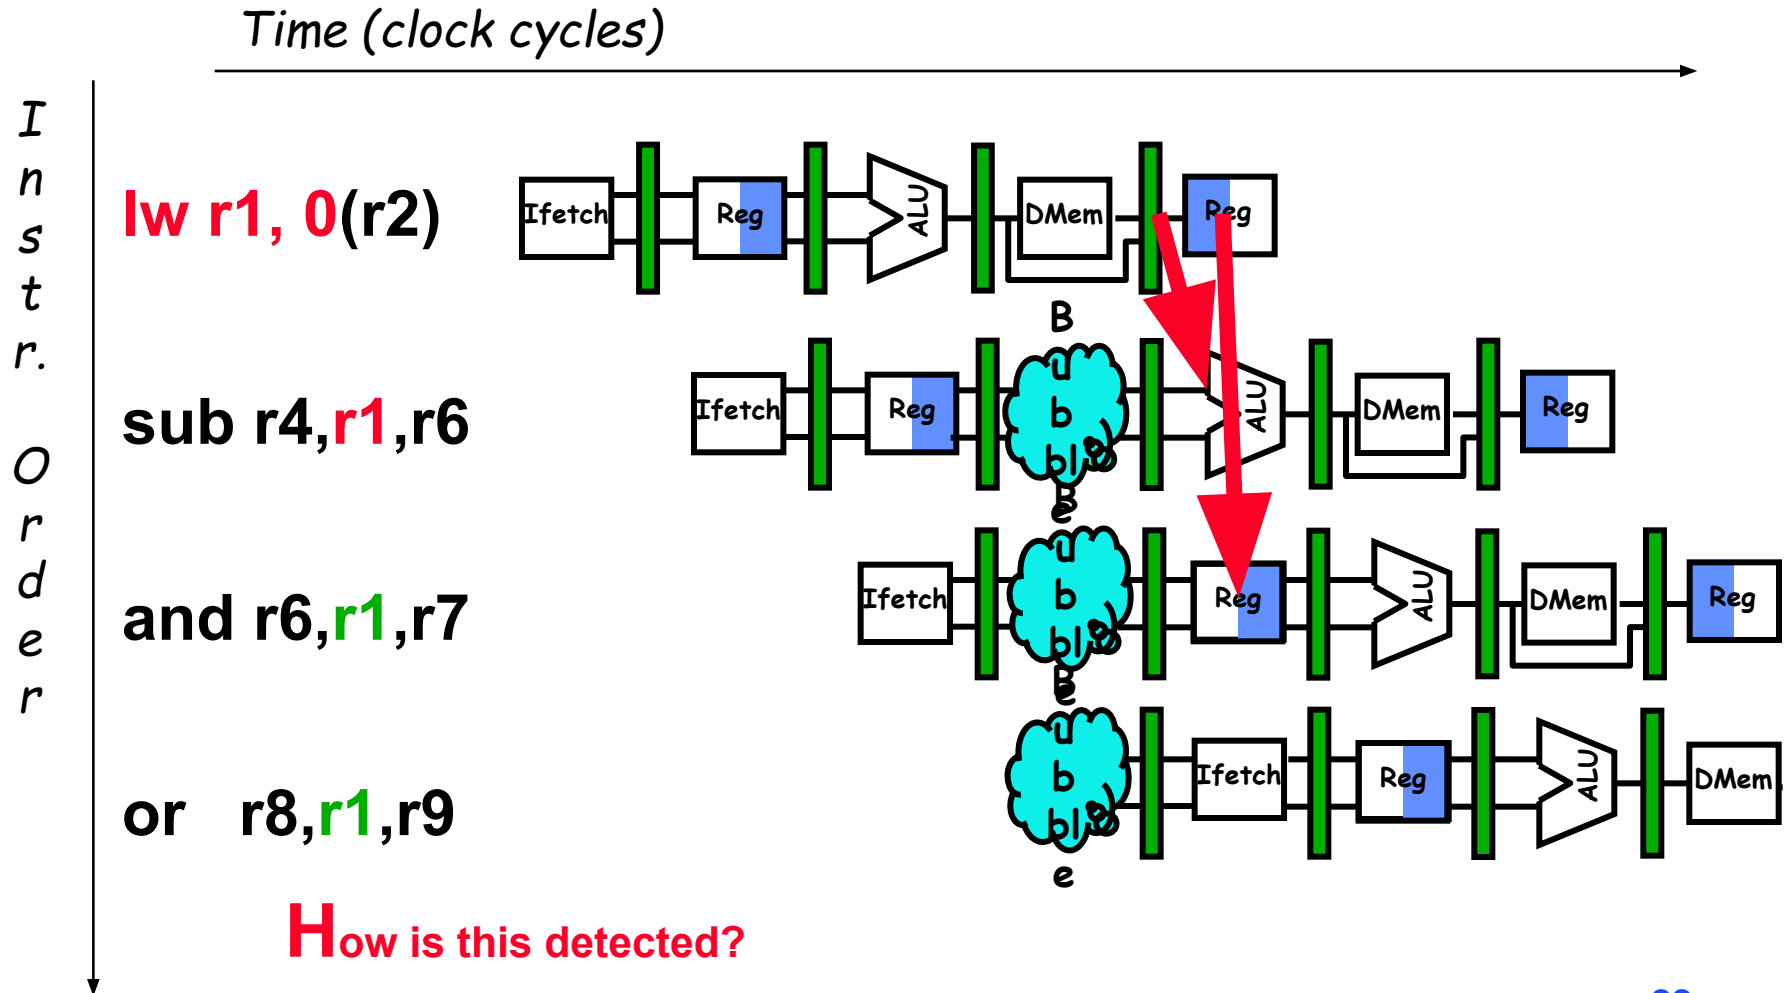
Data Hazard Even with Forwarding

Figure A.9, Page A-21



Data Hazard Even with Forwarding

(Similar to Figure A.10, Page A-21)



Software Scheduling to Avoid Load Hazards

Try producing fast code for

$a = b + c;$

$d = e - f;$

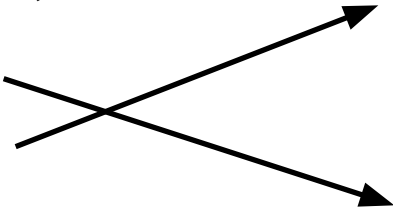
assuming $a, b, c, d, e,$ and f in memory.

Slow code:

```
LW Rb,b
LW Rc,c
ADD Ra,Rb,Rc
SW Ra,a
LW Re,e
LW Rf,f
SUB Rd,Re,Rf
SW Rd,d
```

Fast code:

```
LW Rb,b
LW Rc,c
LW Re,e
ADD Ra,Rb,Rc
LW Rf,f
SW Ra,a
SUB Rd,Re,Rf
SW Rd,d
```



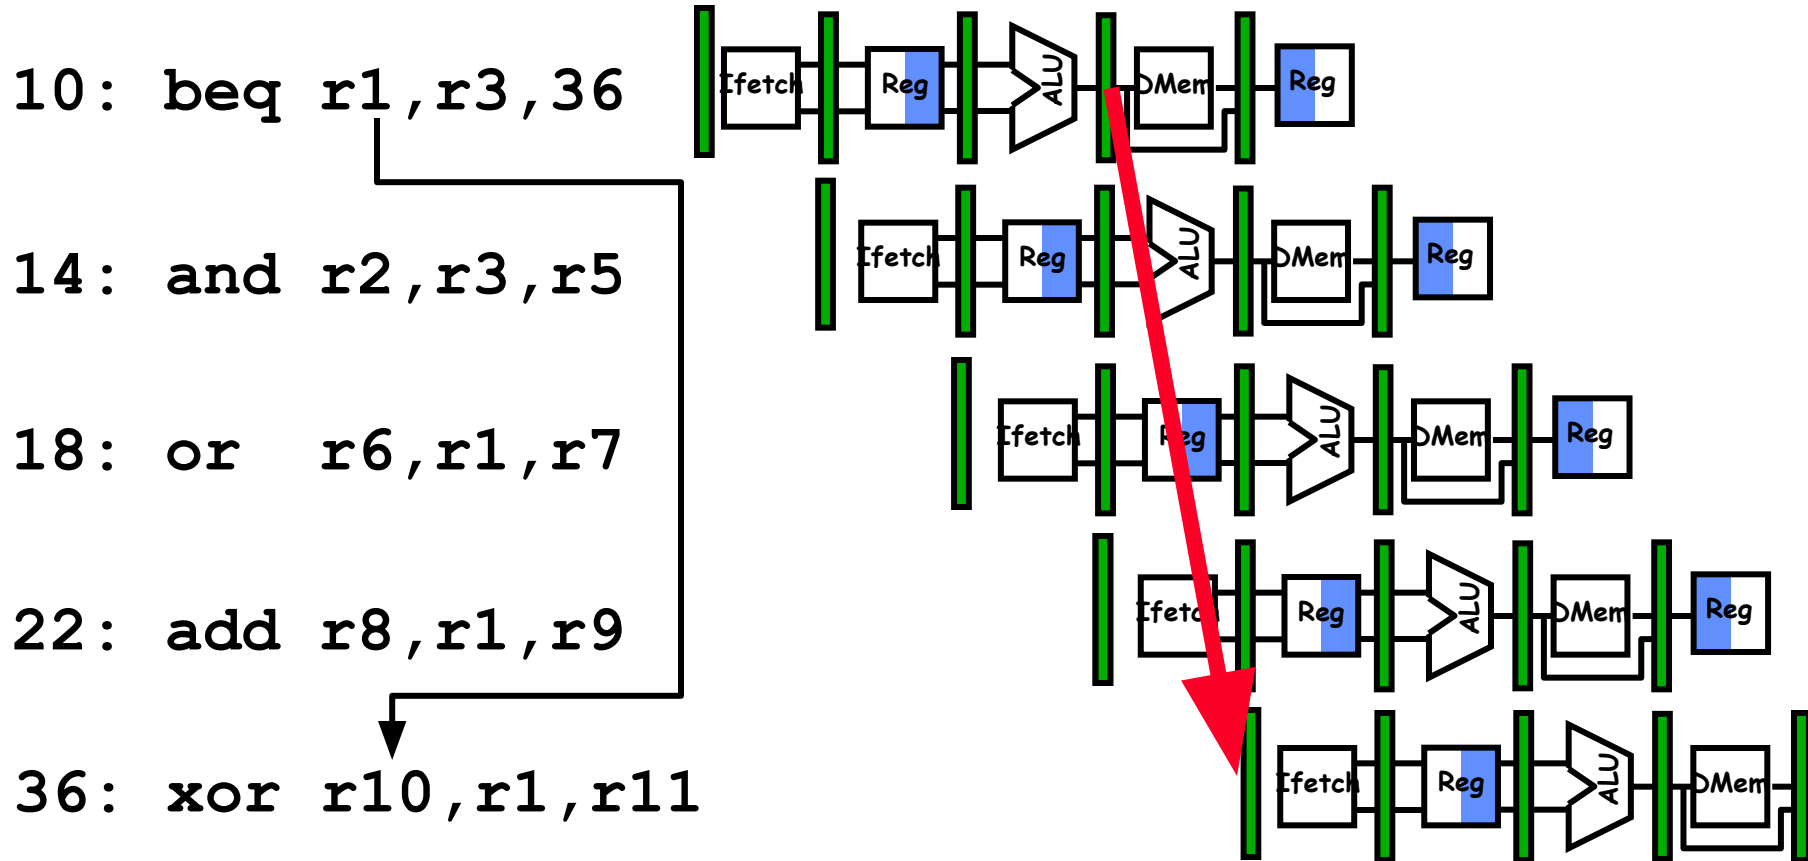
Compiler optimizes for performance. Hardware checks for safety.

Outline

- 5 stage pipelining
- Structural and Data Hazards
- Forwarding
- **Branch Schemes**
- **Multi-cycle Operations**
- **Conclusion**

Control Hazard on Branches

Three Stage Stall



What do you do with the 3 instructions in between?

How do you do it?

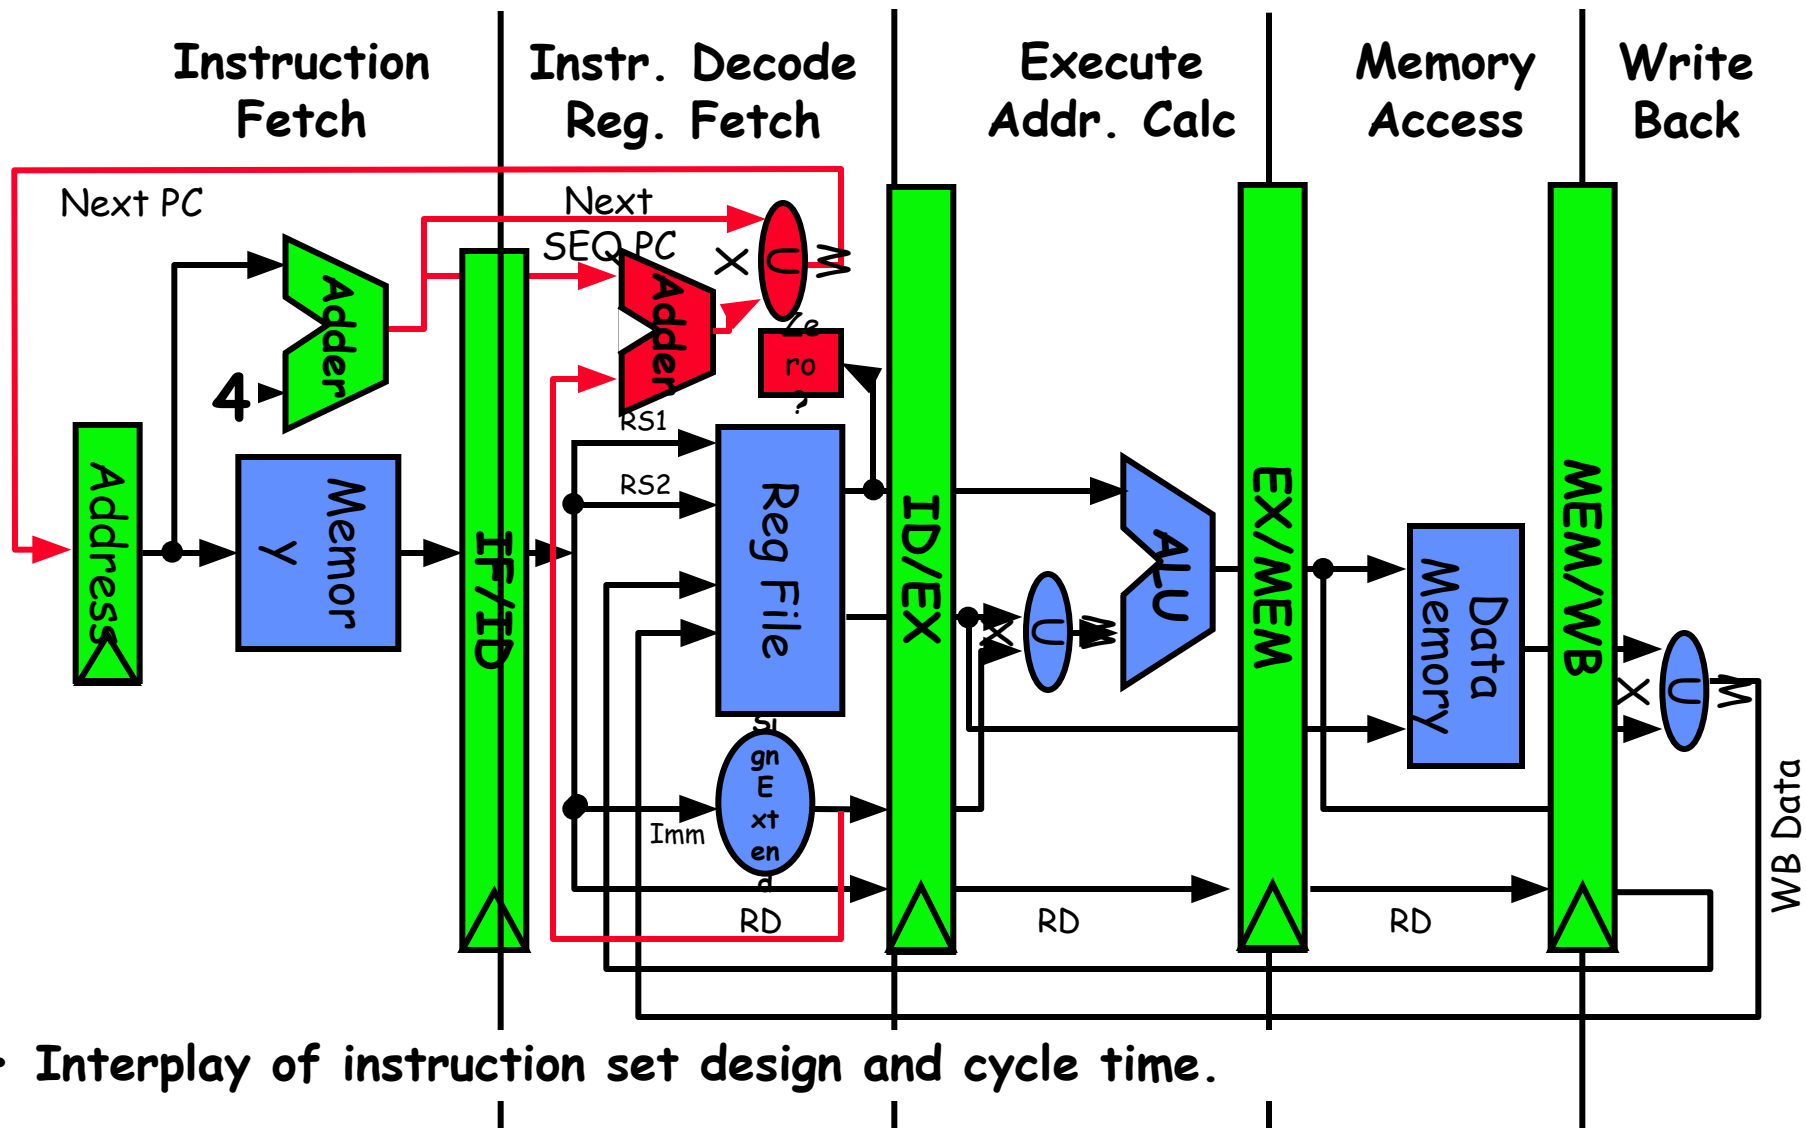
Where is the “commit”?

Branch Stall Impact

- If $\text{CPI} = 1$, 30% branch,
Stall 3 cycles \Rightarrow new $\text{CPI} = 1.9$!
- Two part solution:
 - Determine branch taken or not sooner, AND
 - Compute taken branch address earlier
- MIPS branch tests if register = 0 or $\neq 0$
- MIPS Solution:
 - Move Zero test to ID/RF stage
 - Adder to calculate new PC in ID/RF stage
 - 1 clock cycle penalty for branch versus 3

Pipelined MIPS Datapath

Figure A.24, page A-38



- Interplay of instruction set design and cycle time.

Four Branch Hazard Alternatives

#1: Stall until branch direction is clear

#2: Predict Branch Not Taken

- Execute successor instructions in sequence
- “Squash” instructions in pipeline if branch actually taken
- Advantage of late pipeline state update
- 47% MIPS branches not taken on average
- PC+4 already calculated, so use it to get next instruction


#3: Predict Branch Taken

- 53% MIPS branches taken on average
- But haven't calculated branch target address in MIPS
 - » MIPS still incurs 1 cycle branch penalty
 - » Other machines: branch target known before outcome

Four Branch Hazard Alternatives

#4: Delayed Branch

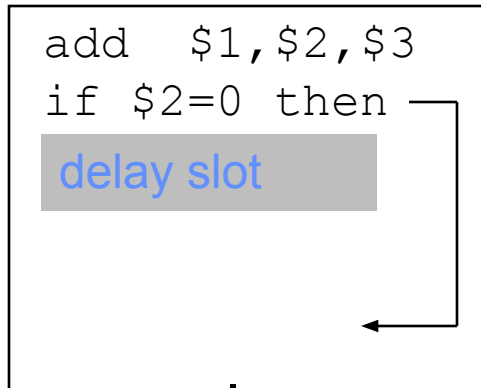
- Define branch to take place **AFTER** a following instruction




- 1 slot delay allows proper decision and branch target address in 5 stage pipeline
- MIPS uses this

Scheduling Branch Delay Slots (Fig A.14)

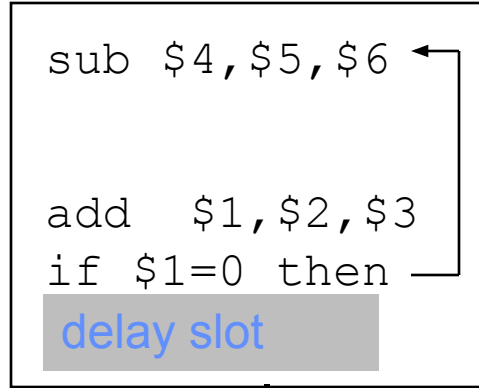
A. From before branch




becomes ↓



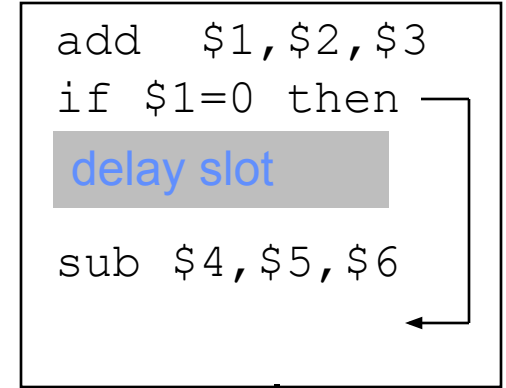
B. From branch target



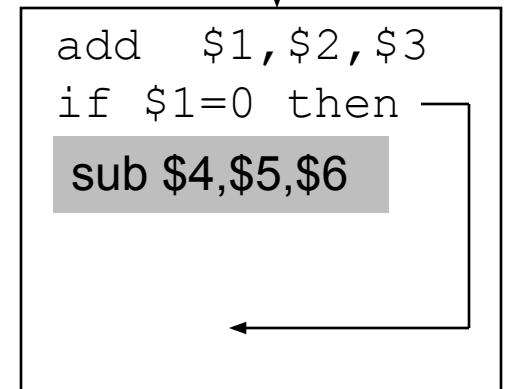
becomes ↓



C. From fall through



becomes ↓



- A is the best choice, fills delay slot & reduces instruction count (IC)
- In B, the `sub` instruction may need to be copied, increasing IC
- In B and C, must be okay to execute `sub` when branch fails

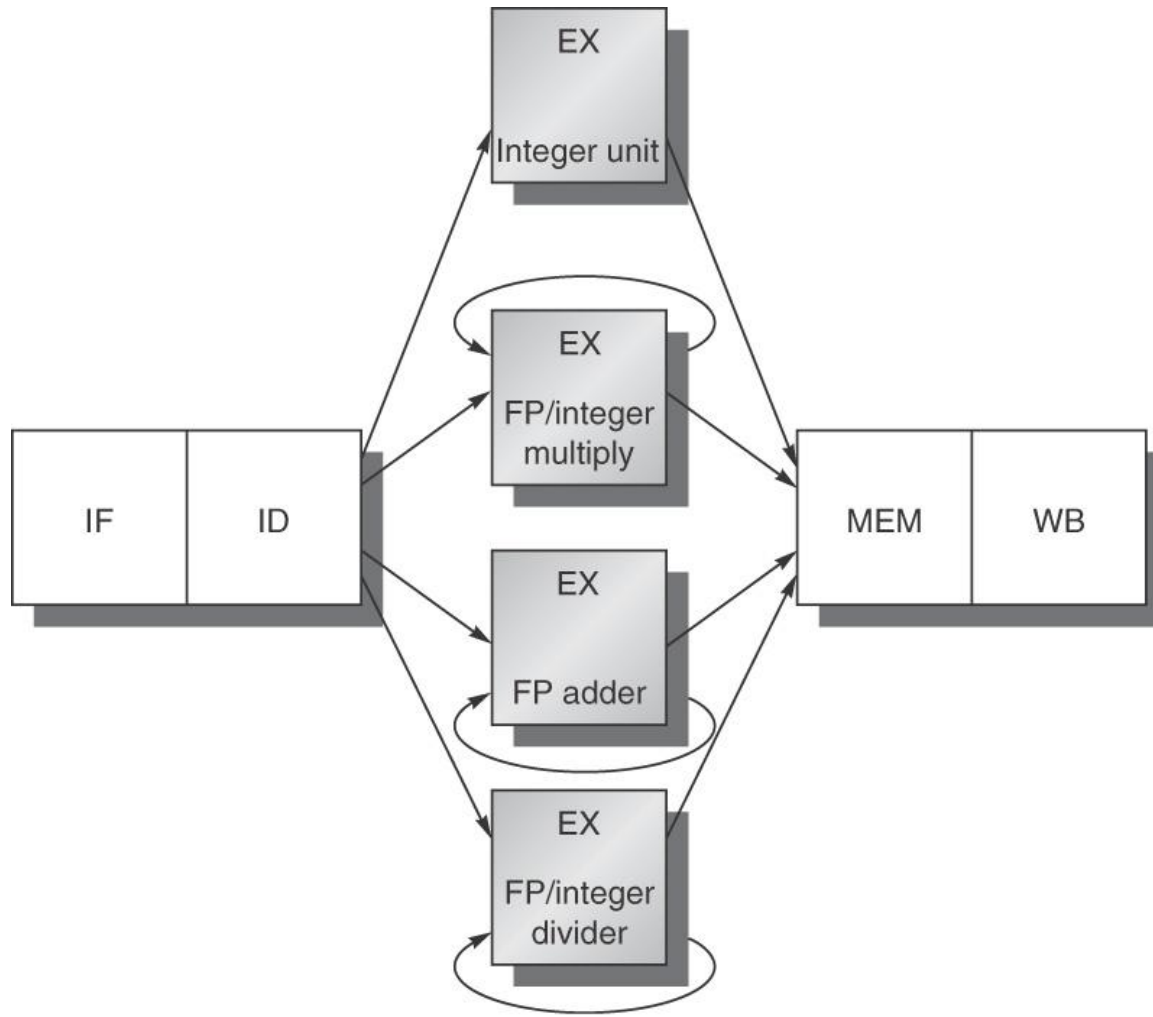
Delayed Branch

- **Compiler effectiveness for single branch delay slot:**
 - Fills about 60% of branch delay slots
 - About 80% of instructions executed in branch delay slots useful in computation
 - About 50% (60% x 80%) of slots usefully filled
- **Delayed Branch downside: As processor go to deeper pipelines and multiple issue, the branch delay grows and need more than one delay slot**
 - Delayed branching has lost popularity compared to more expensive but more flexible dynamic approaches
 - Growth in available transistors has made dynamic approaches relatively cheaper

Outline

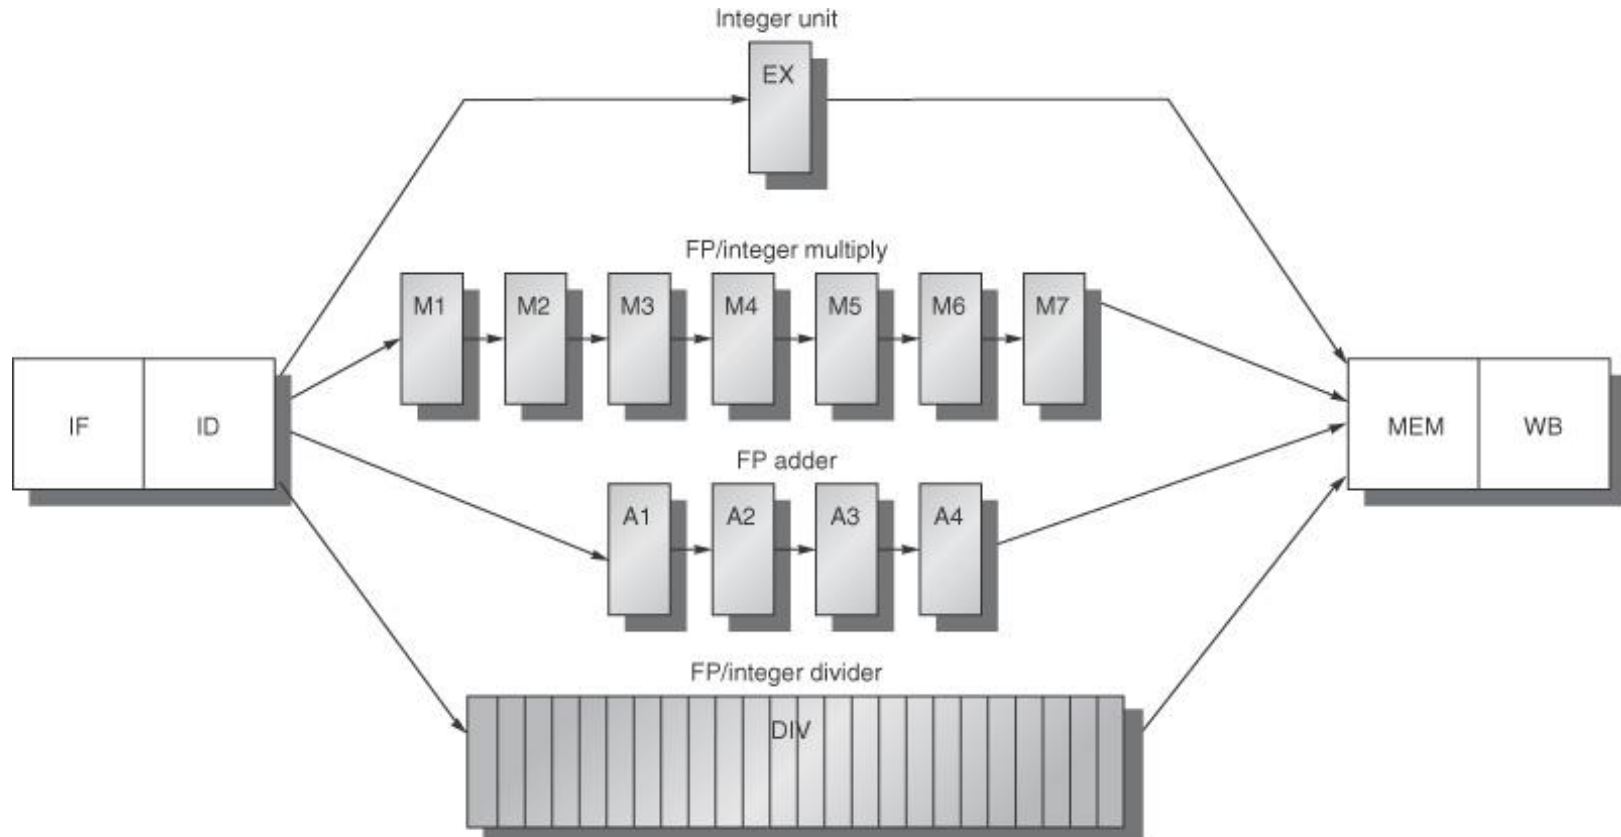
- 5 stage pipelining
- Structural and Data Hazards
- Forwarding
- Branch Schemes
- **Multi-cycle Operations**
- **Conclusion**

Multi-cycle Operations



© 2007 Elsevier, Inc. All rights reserved.

Multi-cycle Operations - Example



© 2007 Elsevier, Inc. All rights reserved.

Multi-cycle Operations - Example

Unit	Latency	Initiation Interval
Integer ALU	0	1
Memory	1	1
FP Adder	3	1
FP Multiplier	6	1
FP Divider	24	25

Problems 1/5

1. Structural hazard on divide

- Solution: Detect at ID stage and stall

Example:

mul.d f0, f4, f6	F	D	M1	M2	M3	M4	M5	M6	M7	M	W
...		F	D	E	M	W					
...			F	D	E	M	W				
add.d f2, f4, f6				F	D	A1	A2	A3	A4	M	W
...					F	D	E	M	W		
...						F	D	E	M	W	
ld.d f2, 0(r2)							F	D	E	M	W

Problems 2/5

2. May have multiple register writes

Solution:

Use shift register to reserve when need to write and stall on conflict

Problems 3/5

3. There are long stalls due to RAW dependencies

Example:

ld.d	f4	, 0(r2)	
mul.d	f0	, f4, f6	1-cycle stall
add.d	f2	, f0, f8	long stall
s.D	f2	, 0(r2)	long stall

Solution:

Mark registers pending as destinations, then in ID stage stall any instruction that uses the marked registers.

Problems 4/5

4. WAW hazards are possible

Example:

```
mul.d  f0, f4, f6
```

```
add.d  f0, f6, f8
```

Solution:

Stall the second instruction if its destination register is pending as a destination.

Problems 5/5

5. Out-of-order completion creates problems in restarting after exceptions

Example:

div.d **f0**, f4, f6 Assume a div by 0 exception
ld.d **f2**, 0(r2) Finishes before div

Solution:

- 1) Accept imprecise exceptions
- 2) Buffer results and write-back in order.

Outline

- 5 stage pipelining
- Structural and Data Hazards
- Forwarding
- Branch Schemes
- Exceptions and Interrupts
- Multi-cycle Operations
- **Conclusion**

And In Conclusion: Control and Pipelining

- **Hazards limit performance on computers:**
 - Structural: need more HW resources
 - Data (RAW,WAR,WAW): need forwarding, compiler scheduling
 - Control: delayed branch, prediction
- **Exceptions, Interrupts add complexity**
- **Multi-cycle operations introduce several problems**

Acknowledgements

- *Adapted from the slides of Prof. David Patterson and Prof. Gheith Abandah.*